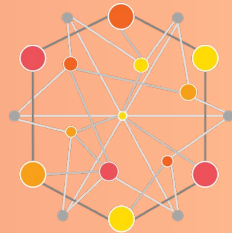


Tutorial: Developing Robust and Scalable Next Generation Workflows Applications and Systems

PEARC 2022



RADICAL Cybertools

`https://radical-cybertools.github.io/`

Introduction

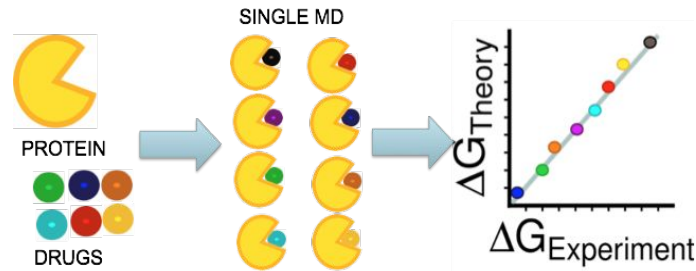
Ensemble-based applications: Applications where the collective outcome of one or more ensembles of tasks is of importance.

Importance of ensemble-based applications:

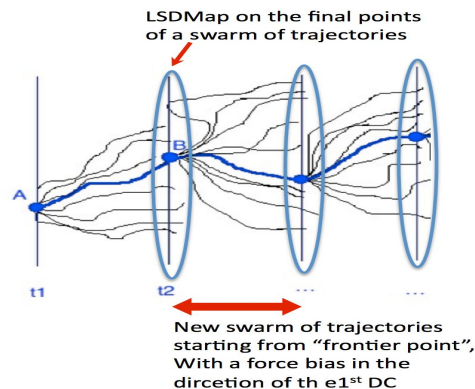
- better algorithms: take advantage of growing **number of processors** on computing infrastructures.

Common in biophysical systems, climate science, seismology, and polar science domains.

RADICAL-Ensemble Toolkit (EnTK): address challenges of **scale, diversity and reliability**



$O(1000)$ simulations to study the binding affinity of multiple drugs on different proteins



$O(1000)$ simulations to understand physical processes of large systems

Advanced Sampling using Adaptive Ensembles

Ensemble-based Adaptive Algorithms:

- runtime data used to determine next stages.



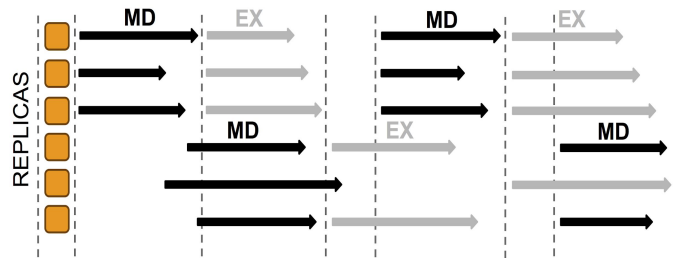
Many biomolecular sampling algorithms formulated as adaptive:

- replica-exchange, Expanded Ensemble, etc.
- Improved simulation efficiency (MSM: 10^3).

Different types of Adaptivity and coupling:

- task parameter(s), order, count, ...



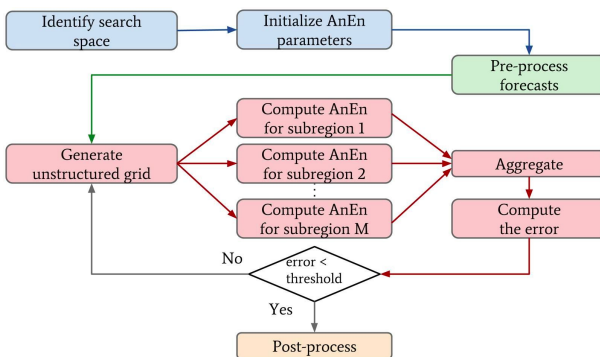
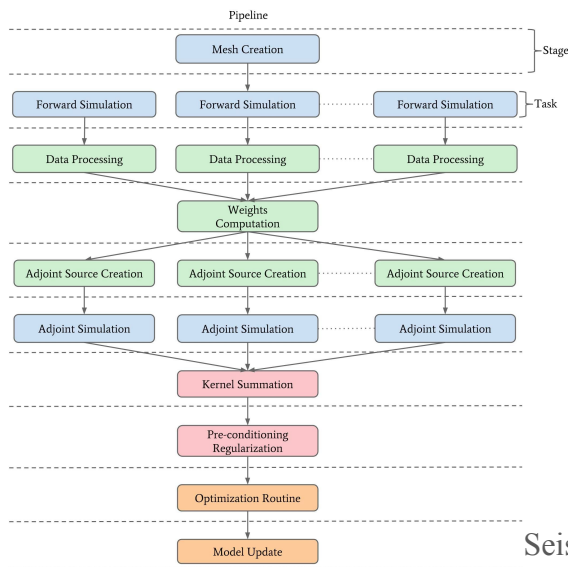
Adaptive logic/MD code should be separated.



Exemplar Use Cases: Seismic Inversion & Meteorological Forecast

Tomographic technique to study the Earth's interior.

Probabilistic meteorological forecast to study optimal placement of solar power plants.



RADICAL-EnTK: API Abstractions

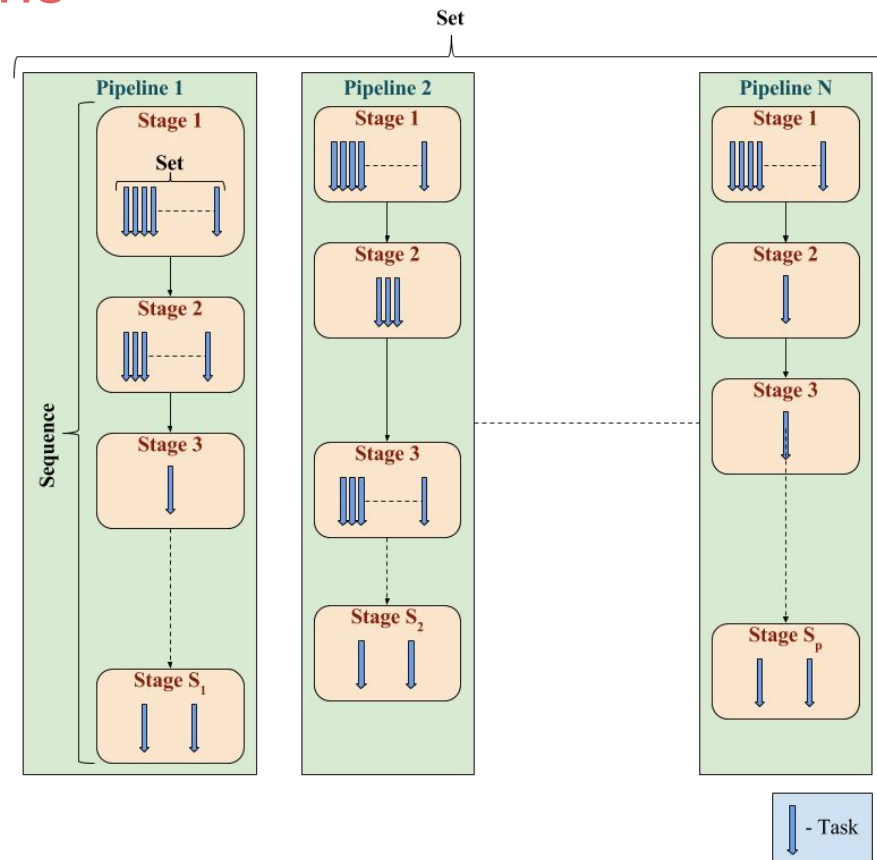
Two (pythonic) collections of objects:

- **Set:** contains objects that have no relative order with each other
- **Sequence/List:** contains objects that have a linear order, i.e. object 'i' depends on object 'i-1'

Three objects:

- **Task:** description of executing kernel
- **Stage:** set of Tasks, i.e. all tasks of a stage may execute concurrently
- **Pipeline:** sequence of Stages, i.e. Stage 2 may only commence after Stage 1 completes

*An ensemble **application** is described as a set of Pipelines.*



RADICAL-EnTK Example 1: Ensemble of Simulation Pipelines

Each simulation pipeline has multiple stages

Simulation stage consists of multiple tasks

Data dependencies exist between stages

Prepare tutorial environment

```
cd $HOME  
cd tutorial/  
. 1-ensemble-rct.env  
cd 1-ensemble-rct
```

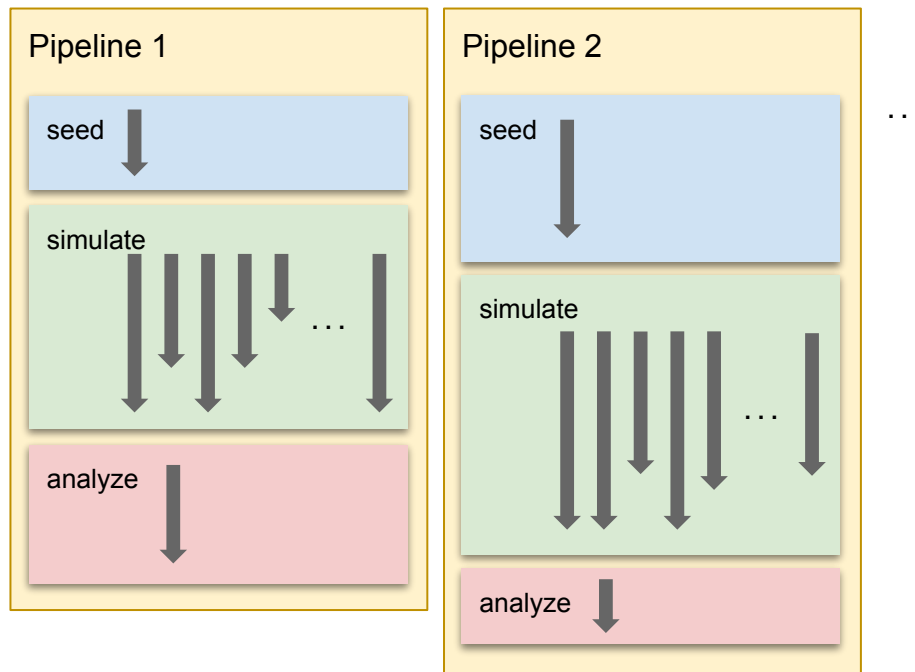
Open first example

```
vim radical_entk_1.py
```

Run example

```
./radical_entk_1.py
```

Original files are backed up in `orig/`



RADICAL-EnTK Example 1: Exercises

1. Switch execution backend to use FLUX instead of default `fork/exec`^{*}:
 - Look at the `appman.resource_desc` in the program's `main` section
 - Solution is in `solution_1.1.py`
2. Change the number of ensemble members (number of pipelines) and number of simulations per pipeline:
 - You are presumably running on a small resource - be gentle ;-)
 - Look at the `for` loop in the program's `main` section
 - Look at the construction of Stage 2 (`s2`)
 - Solution is in `solution_1.2.py`
3. Add a fourth stage which computes the square root of the sum:
 - The kernel could be something like: `echo "sqrt($(cat sum.txt))" | bc`
 - Output staging should move from previous last stage (`s3`) to the new stage (`s4`)
 - Solution is in `solution_1.3.py`

(*) RCT execution backends include Slurm, LSF, APRUN, PRTE, mpirun/mpiexec, ibrun, etc.

RADICAL-EnTK Example 2: Ensemble of Dynamic Pipelines

After each simulation stage, check results

if divergent: re-seed the pipeline

elif convergent: finish pipeline

else: continue simulation

Stage `'post_exec'` directives are used to decide on pipeline progression:

- Intermediate data are staged out
- New stages are added dynamically
- Pipelines are completed as needed

Reset tutorial environment

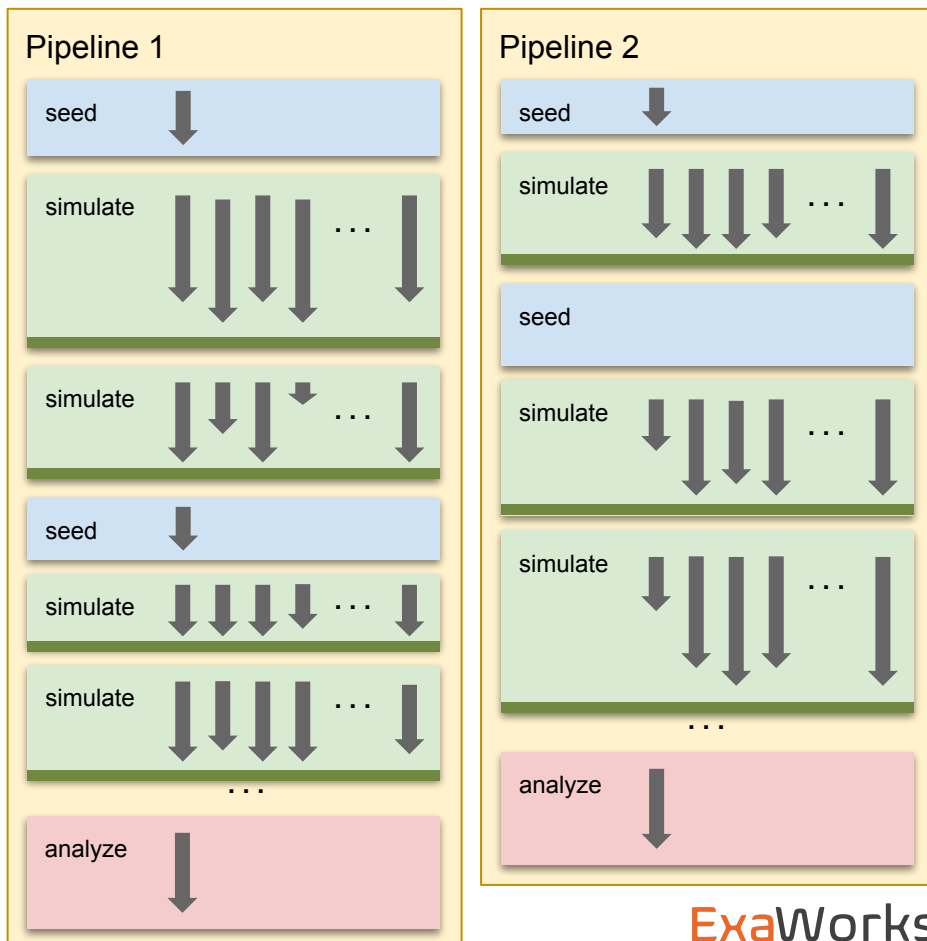
```
make clean
```

Open second example

```
vim radical_entk_2.py
```

Run second example

```
./radical_entk_2.py
```



RADICAL-EnTK Example 2: Exercises

Calculating intermediate results is costly: all data need to be staged back and analyzed. Instead, insert a **`stage 2b`** which analyzes data on the target resource and only then stages back the result to decide about pipeline continuation.

- This exercise is very similar to previous exercise 3
- **`stage 2b`** would be very similar to **`stage 3`**
- Consider what stage then needs to hold the **`post_exec`**
- Solution is in **`solution_2.1.py`**
- Note: to avoid very long output, you can turn off the reporter with:

```
export RADICAL_REPORT=False
```

Exercises

`https://tinyurl.com/exaworks`