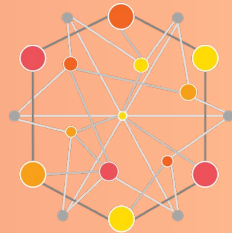


# Tutorial: Developing Robust and Scalable Next Generation Workflows Applications and Systems

*PEARC 2022*



# Swift/T

---

`http://swift-lang.org/Swift-T`

# Swift/T: Enabling high-performance scripted workflows

Write site-independent scripts, translates to MPI

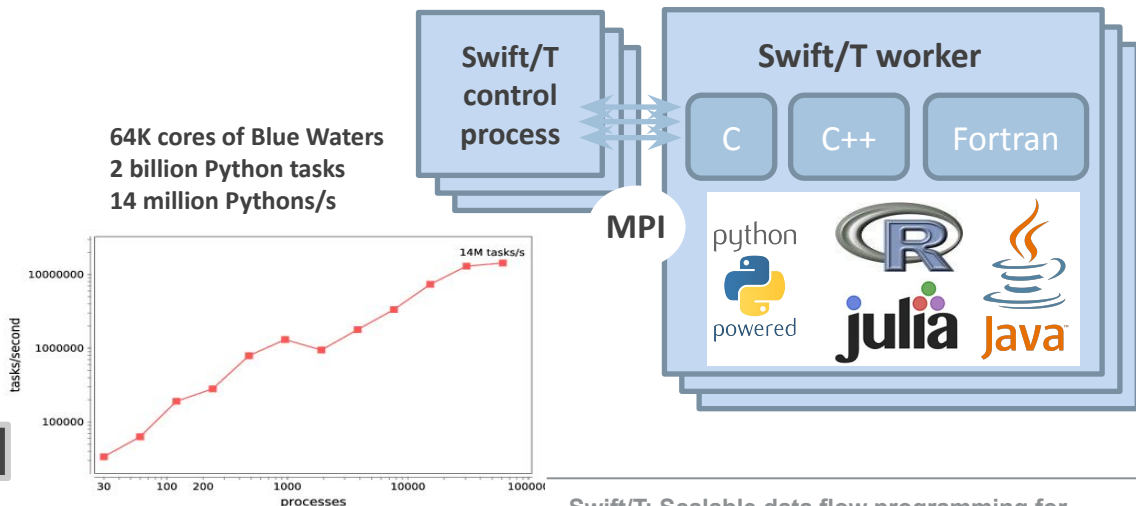
Automatic task parallelization and data movement

Invoke native code, script fragments  
in Python and R

Rapidly subdivide large  
partitions for MPI jobs  
in multiple ways (MPI 3.0)

```
$ spack install stc
```

```
$ conda install -c lightsource2-tag swift-t
```



Swift/T: Scalable data flow programming for  
distributed-memory task-parallel applications  
Proc. CCGrid 2013.

# Goals of the Swift language

Swift was designed to handle many aspects of the computing campaign

- Make it easy to run large batteries of external program or library executions
- Ability to integrate many application components into a new workflow application
- Enable complex tasks based in other scripting languages (e.g., Python) or parallel MPI tasks
- Provide rich programming language at the top level – fully generic
- Data structures for complex data organization
- Portability- separate site-specific configuration from application logic
- Logging, provenance, and plotting features
- Support implicit concurrency and conventional programming constructs

# The Swift programming model

## All progress driven by concurrent dataflow

```
(int r) myproc (int i, int j)
{
    int x = F(i);
    int y = G(j);
    r = x + y;
}
```

- `F()` and `G()` implemented in native code or external programs
- `F()` and `G()` run concurrently in different processes
- `r` is computed when they are both done
- This parallelism is *automatic*
- Works recursively throughout the program's call graph

# Swift syntax

- Data types

```
int i = 4;
string s = "hello world";
file image<"snapshot.jpg">;
```

- Structured data

```
typedef image file;
image A[];
type protein_run {
    file pdb_in; file sim_out;
}
bag<blob>[] B;
```

- Conventional expressions

```
if (x == 3) {
    y = x+2;
    s = strcat("y: ", y);
}
```

- Parallel loops

```
foreach f,i in A {
    B[i] = convert(A[i]);
}
```

- Data flow

```
merge(analyze(B[0], B[1]),
      analyze(B[2], B[3]));
```

- 
- **Swift: A language for distributed parallel scripting.** J. Parallel Computing, 2011
  - **Compiler techniques for massively scalable implicit task parallelism.** Proc. SC, 2014

# Swift task invocation

- Shell access

```
app (file o) myapp(file f, int i)
{"/mysim" "-s" i @f @o; }
```

- Or simply invoke a string:

```
output,error = system1("echo HELLO");
```

or

```
output,error = system(["echo", "HELLO"]);
```

- Python and R (etc.)

- ```
python("import numpy as np",
      "repr(np.eye(1))");
```

- ```
R("A=c(4,5,6)", "toString(mean(A))");
```

- External MPI programs

```
@par=8 launch("./my-program", args[], envs[]);
```

- In-memory MPI libraries

```
@par (string z)
covid_model_run(string config, string params)
"covid_model" "0.0" "covid_model_tcl";

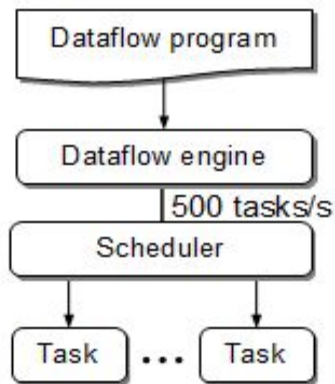
@par=procs_per_run covid_model_run(default_model_props, p);
```

- mpi4py

```
@par=6 python_parallel_persist(
    "import test_6;test_6.f('HELLO')");
test_6.py:f():
comm = turbine_helpers.get_task_comm()
size = comm.Get_size() ; rank = comm.Get_rank()
comm.barrier()
...
```

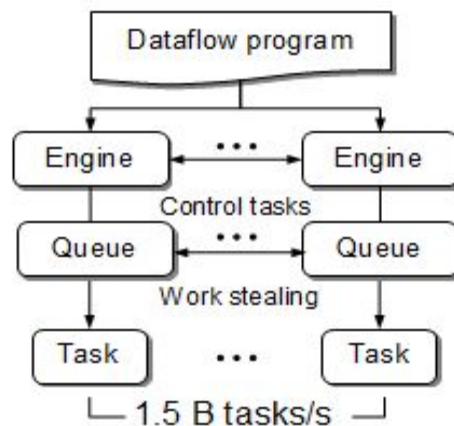
# Centralized evaluation is a bottleneck at extreme scales

Had this (Swift/K):



Centralized evaluation

Now have this (Swift/T):



Distributed evaluation

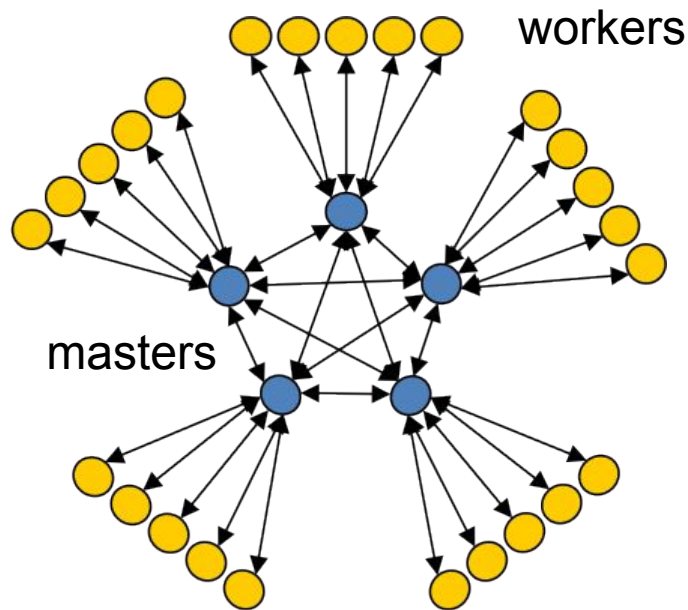
**Turbine: A distributed-memory dataflow engine for high performance many-task applications.** Fundamenta Informaticae 28(3), 2013



# Asynchronous Dynamic Load Balancer

## ADLB for short

- An MPI library for master-worker workloads in C
- Uses a variable-size, scalable network of servers
- Servers implement work-stealing
- The work unit is a byte array
- Optional work priorities, targets, types
- For Swift/T, we added:
  - Server-stored data
  - Data-dependent execution

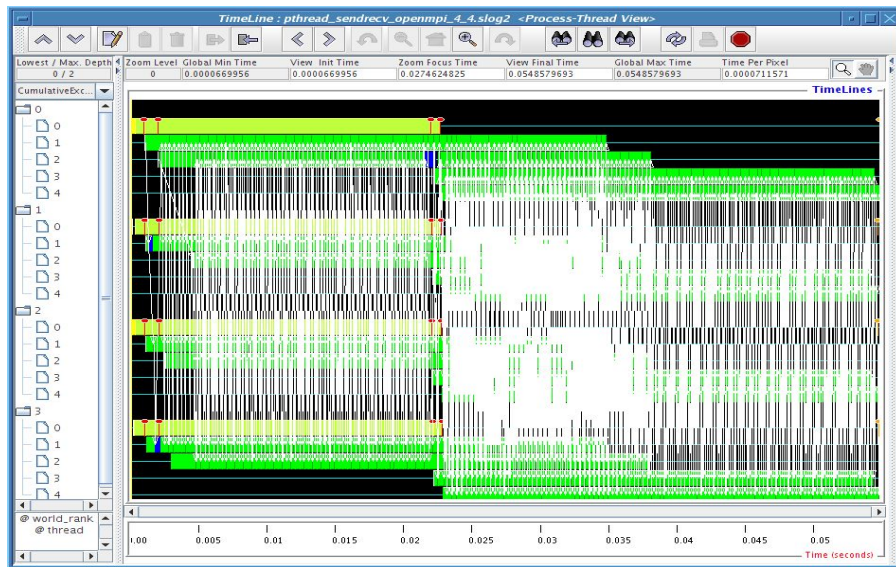


- 
- Lusk et al. **More scalability, less pain: A simple programming model and its implementation for extreme computing.** SciDAC Review 17, 2010

# MPI: The message passing interface

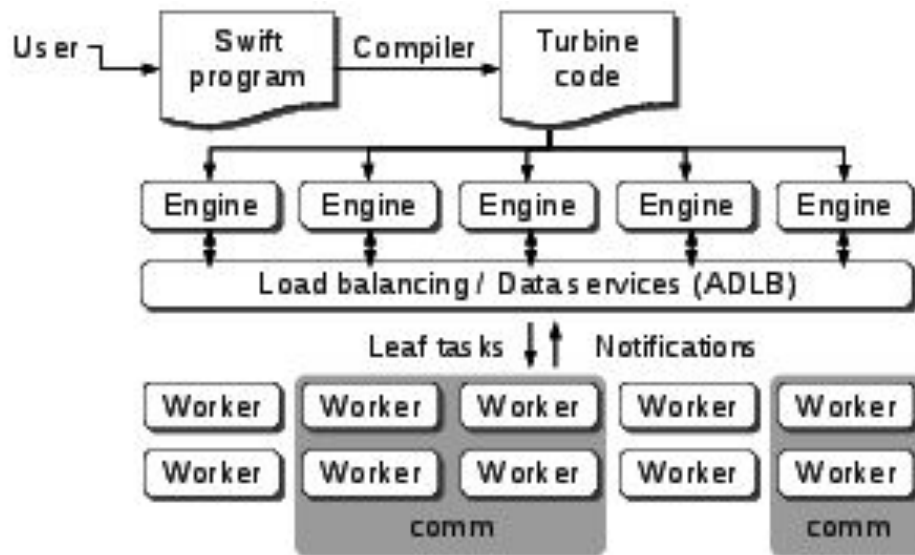


- Programming model used on large supercomputers
- Can run on many networks, including sockets, or shared memory
- Standard API for C and Fortran; other languages have working implementations
- Contains communication calls for
  - Point-to-point (send/recv)
  - Collectives (broadcast, reduce, etc.)
- Interesting concepts
  - Communicators: collections of communicating processing and a context
  - Data types: Language-independent data marshaling scheme



# Parallel tasks in Swift

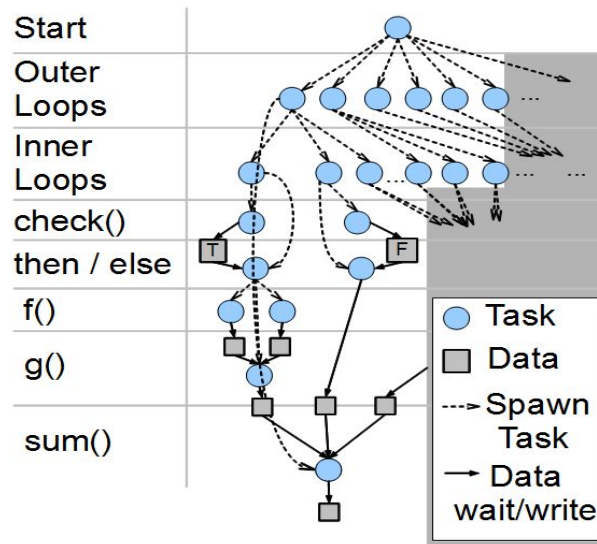
- Swift expression: `z = @par=8 f(x,y);`
- When `x, y` are stored, Turbine releases task `f` with `parallelism=8`
- Performs `ADLB_Put(f, parallelism=8)`
- Each worker performs `ADLB_Get(&task, &comm)`
- ADLB server finds 8 available workers
- Workers receive ranks from server
  - Perform `MPI_Comm_create_group()`
- `ADLB_Get()` returns:  
`task=f, size(comm)=8`
- Workers perform user task
  - communicate on `comm`
- `comm` is released by Turbine
- Can hand the communicator to RepastHPC, LAMMPS, NAMD, DIY, CODES/ROSS, etc.



- Wozniak et al. Dataflow coordination of data-parallel tasks via MPI 3.0. Proc EuroMPI, 2013.

# Swift/T: Fully parallel evaluation of complex scripts

```
int X = 100, Y = 100;
int A[][];
int B[];
foreach x in [0:X-1] {
  foreach y in [0:Y-1] {
    if (check(x, y)) {
      A[x][y] = g(f(x), f(y));
    } else {
      A[x][y] = 0;
    }
  }
  B[x] = sum(A[x]);
}
```



Compiler techniques for massively scalable implicit task parallelism.  
SC 2014.

# City-scale COVID-19 epidemic modeling



ACM Gordon Bell Special Prize  
for High Performance Computing-Based  
COVID-19 Research  
Finalist  
2020

- **Observed city data**

- Hospitalizations, cumulative deaths from Chicago Department of Public Health
- Detailed line list data from Illinois Department of Public Health

- **ML Phases:**

- Distribution of R code snippets via R foreach %dopar% syntax; work distributed by EMEWS

- **Simulation:**

- Distribution of MPI tasks via Swift/T @par syntax.  
MPI communicators are dynamically allocated over in-order cores by Swift/T using MPI\_Comm\_create\_group()

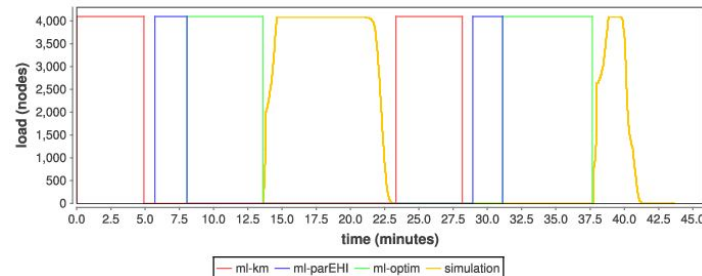
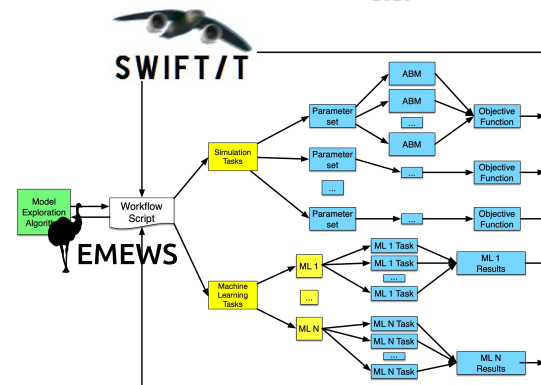
- **Key performance metric:** Scalability and Time to Solution

- Keep the cores busy in presence of changing task types and workflow dynamics

- **Simulation phase** starts 1024 MPI tasks, each on a 256-rank MPI communicator

- Assigns tasks at rate of 4,695 core-tasks per second,  
73 communicators per second

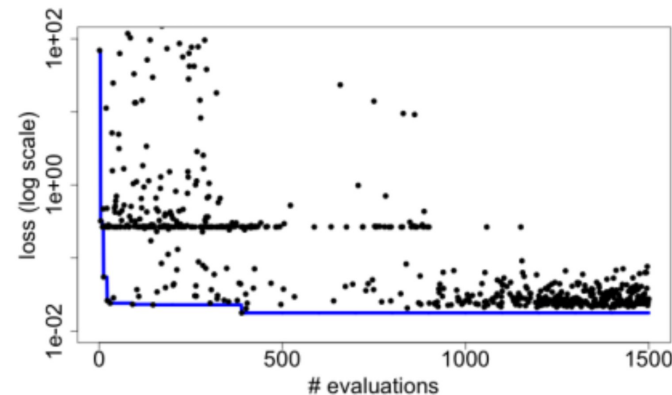
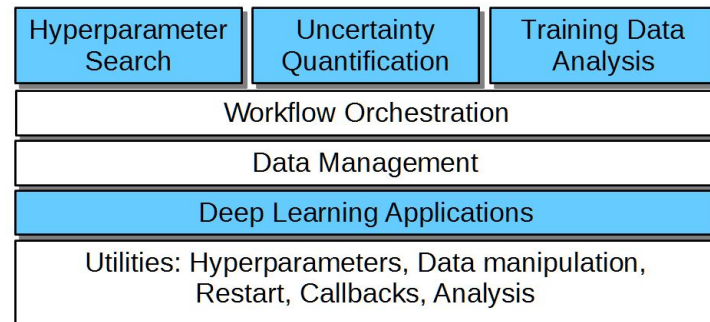
- **The ML phases** are single-node, vendor-optimized R calls



A population data-driven workflow for COVID-19 modeling  
and learning. IJHPCA 2021.

# ECP CANDLE Hyperparameter optimization

- Goal: Develop an exascale deep learning environment for cancer, enabling the most challenging deep learning problems in cancer research to run on the most capable supercomputers
  - Neural networks have a large number of possible configuration parameters, called hyperparameters
  - CANDLE/Supervisor consists of several high-level workflows
  - Capable of modifying/controlling application parameters dynamically as the workflow progresses and training runs complete
  - Distribute work across large computing infrastructure, manage progress
  - Underlying applications are Python programs that use Keras/TensorFlow
- Hyperparameter search plot:
  - Search trajectory of mlrMBO (R model-based optimization) algorithm
  - Each iteration does 300 evaluations (batch size)
  - Minimum and average performance on validation data set decreases as the ME algorithm learns

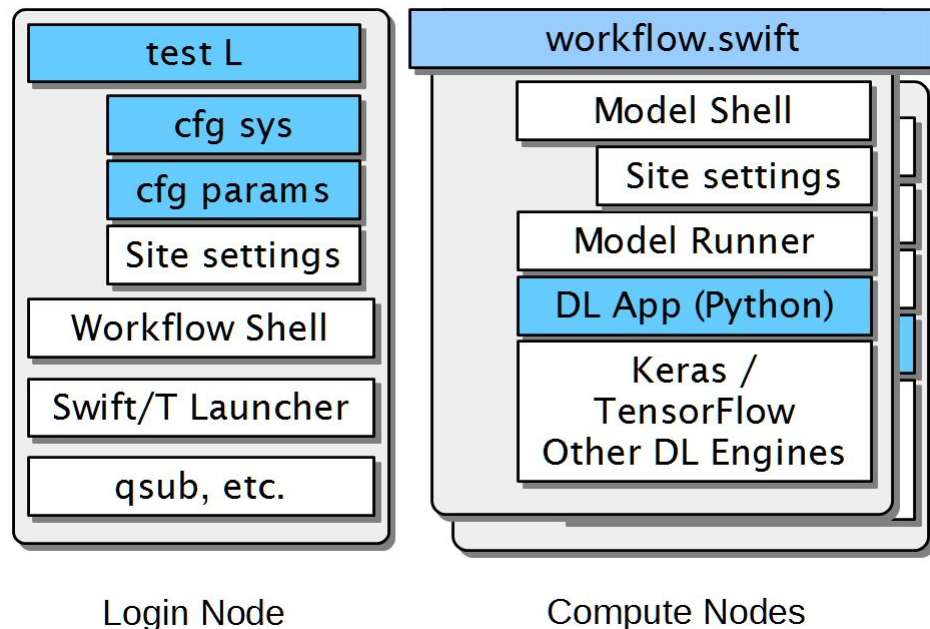


CANDLE/Supervisor: A workflow framework for machine learning applied to cancer research. BMC Bioinform. 2018.

# CANDLE/Supervisor Implementation

## Script schematic

- Runs start with a test script
- CFG scripts contain settings for a system or parameters for a given study (e.g., search space)
- Reusable site settings
- The workflow shell script sets up the run
- Swift/T launches and manages the workflow
- Reusable Model scripts set up each app run
- The DL app uses Keras/TF plus CANDLE Python utilities



# Exercises

---

`https://tinyurl.com/exaworks`



# How to run the tutorial exercises

- In the provided instance,

1. `source ~/tutorial/2-workflow-dl-swift.env`
2. `cd tutorial/2-workflow-dl-swift/2-workflow-dl-swift`
3. `git pull`



# Installation

---

`http://swift-lang.github.io/swift-t/guide.html#install\_source`

# New MPI\_LAUNCH feature

- Allows Swift/T to run external parallel programs on subcommunicators inside a large allocation on a big machine

- Swift/T syntax:

```
@par=8 launch("./my-program", args[], envs[]);
```

- Provides:
  - Scalable, in-place job launch
  - Handles cases where called program crashes
  - Can pack many such variably-sized programs within a large workflow

# Newer MPI\_LAUNCH\_MULTI feature

- Allows Swift/T to run external parallel program groups on subcommunicators inside a large allocation on a big machine
  - Call these groups Functional Online Bundles (FOBs)
- Swift/T syntax:  

```
@par=8 launch_multi(procs[], programs[], args[][] , envs[][] ,  
                    <colors>);
```
- Provides:
  - Scalable, in-place simultaneous job launch
  - The programs are able to find each other and communicate with ADIOS (or other techniques)
  - Job layout can be controlled with the optional colors argument
  - A variety of other controls are available via special environment variables

# Swift/T example

**When file A is created, launch N sub jobs of varying size**

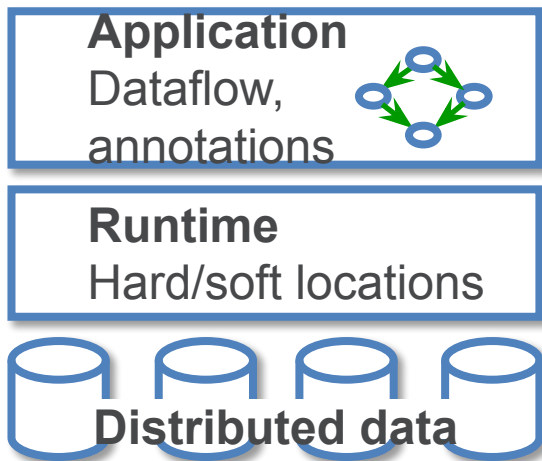
```
file B[]; // Define an array of file variables
A => {
  foreach i in [0:N-1] {
    file B_i<"B-%i.txt"%i>;
    string args_B[] = [ int2string(i),
                        filename(A), filename(B_i) ];
    @par=i launch("./child.x", args_B) => B_i = touch();
    B[i] = B_i;
  }}
}
```

- Child tasks are load-balanced, MPI\_Comm\_create\_group() is done automatically!

# Features for Big Data Analysis

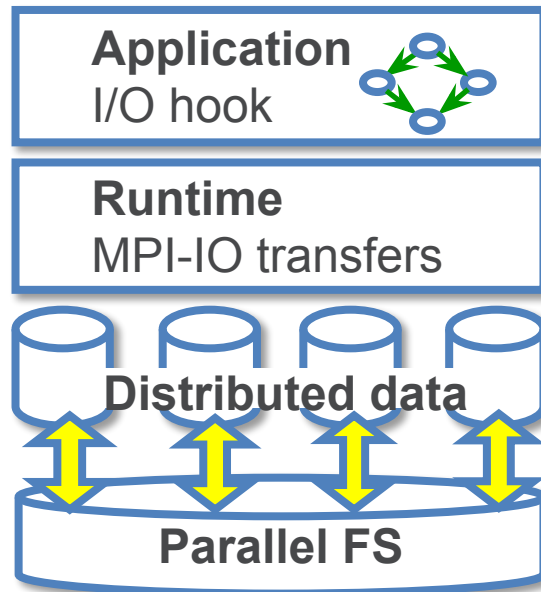
- **Location-aware scheduling**

User and runtime coordinate data/task locations



- **Collective I/O**

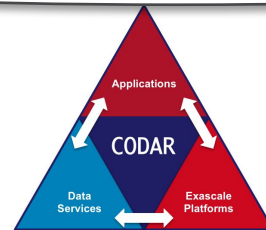
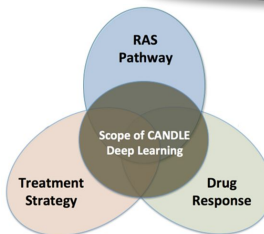
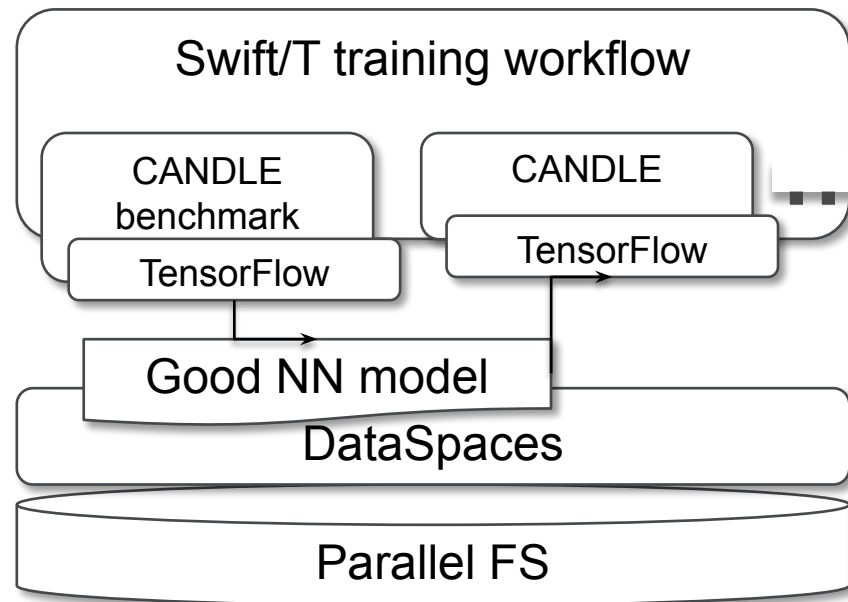
User and runtime coordinate data/task locations



- **Big data staging with MPI-IO for interactive X-ray science.** Wozniak et al. Proc. Big Data Computing 2014.
- **Experimental evaluation of a flexible I/O architecture for accelerating workflow engines in ultrascale environments.** F. Duro, Wozniak, et al. Parallel Computing 61, 2017.

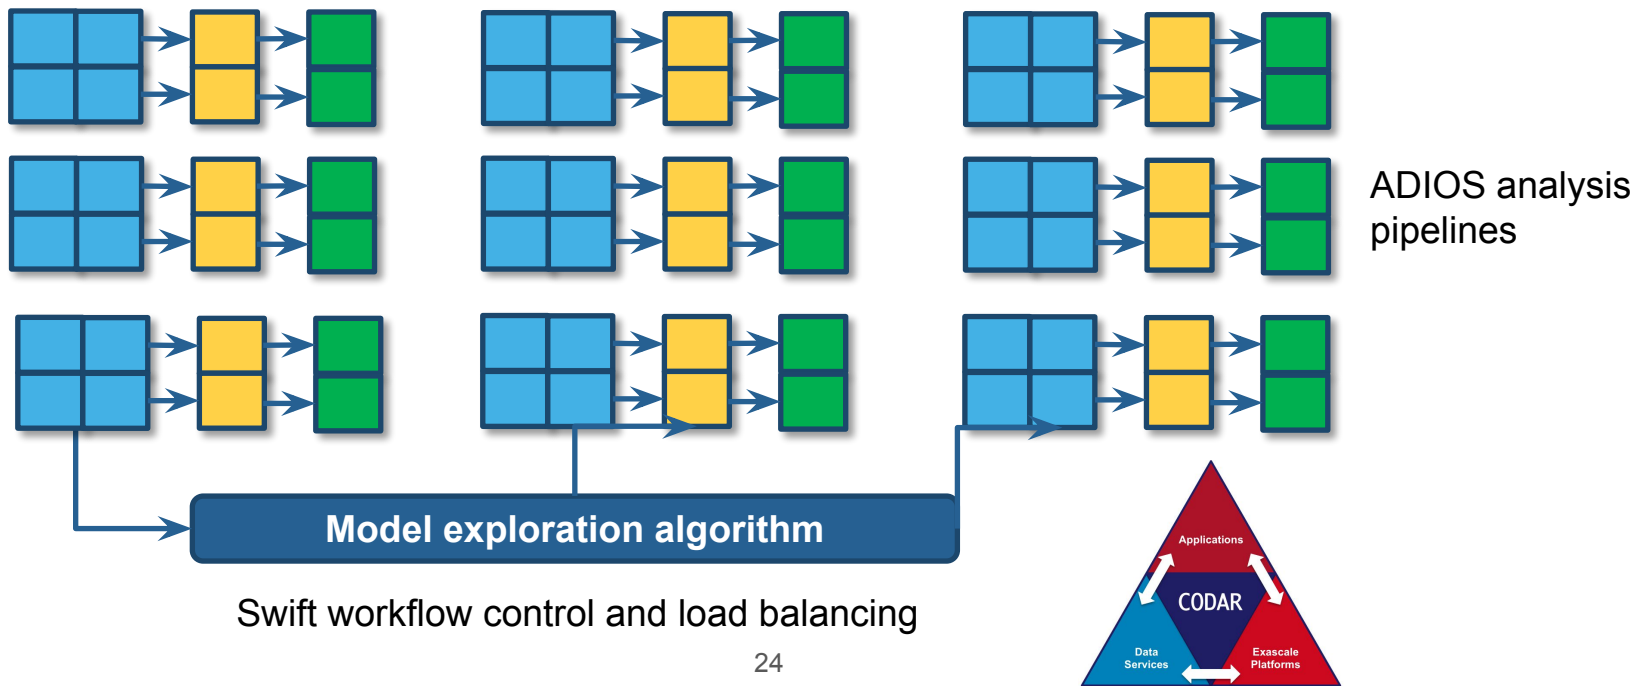
# ECP INTERACTION: CODAR, CANDLE

- **CANDLE** workflows produce a great number of medium-sized ML models
- **Goal:** Cache these on compute node storage for *possible* later use. Need to flush to global FS before end of run, but many models will be discarded
- **Approach:** Integrated Swift/T workflow system used in CANDLE with DataSpaces client
- Provide an opportunity for workflow-based data analysis and I/O reduction
- Demonstrate the utility of node-local storage for complex workflows
- **Scaling deep learning for cancer with advanced workflow storage integration.**  
Proc. MLHPC @ SC 2018.



# ECP CODAR: Workflows of ADIOS transfers

- Enable Swift to dynamically lay out tasks
- Control large simulation/redistribute/analysis ensembles
- Highly flexible, programmable use of MPI subjobs





# U. Chicago Hospital: Cancer ensembles

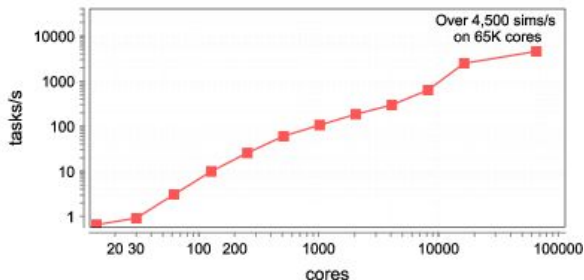
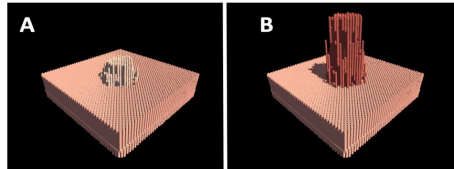
## Best paper at SC Cancer Workshop 2016

- Parameter fitting for biological phenomenon (DNA repair rate) via massive scale evolutionary algorithm in Swift/T framework



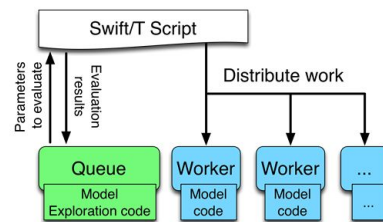
### GIOABM – Integration into SEGMEiT

- A cancerous cell has three features: immortality, invasiveness, and ability to proliferate unnaturally
- GIOABM call functionality overlaps with SEGMEiT at four locations:
  - B-catenin: proliferation
  - PI3K: Proliferation/Apoptosis
  - TGF- $\beta$ /SMAD: Proliferation/Apoptosis
  - P53: Gene repair/Apoptosis
- Added E-cadherin protein mutation to SEGMEiT representing invasiveness



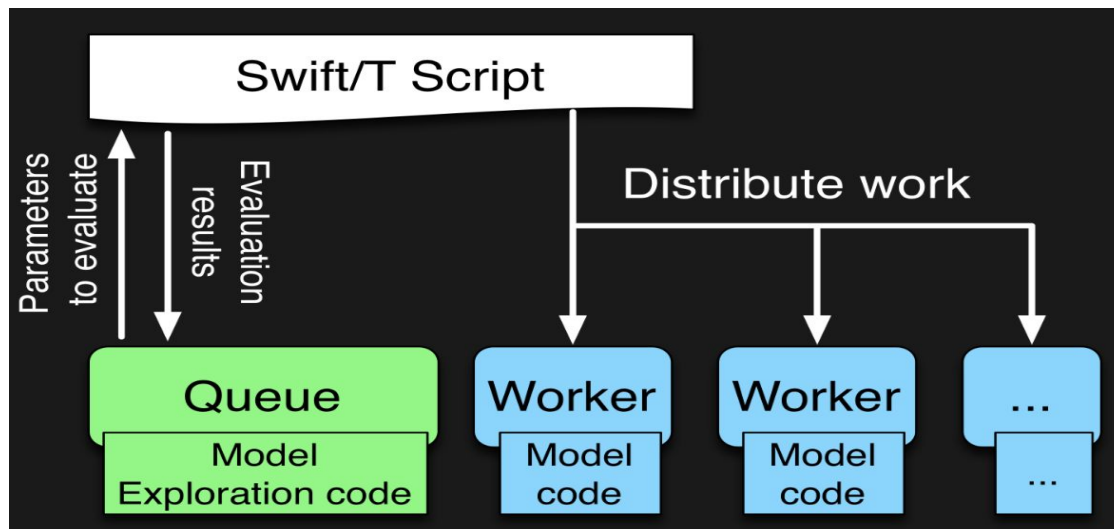
### Extreme-scale Model Exploration with Swift (EMEWS)

- EMEWS offers:
  - the capability to run very **large**, highly **concurrent** ensembles of simulations of **varying types**
  - supports a **wide class of ME algorithms**, including those increasingly available to the community via Python and R libraries
- EMEWS design goal: to ease software integration while providing scalability to the largest scale (petascale plus) supercomputers, running millions of models

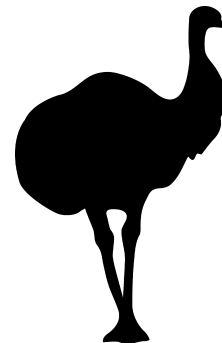


- Anatomic-scale cancer modeling using the Extreme-scale Model Exploration with Swift (EMEWS) framework. Proc. Cancer Workshop @ SC, 2016. (Best paper)

# EMEWS workflow structure

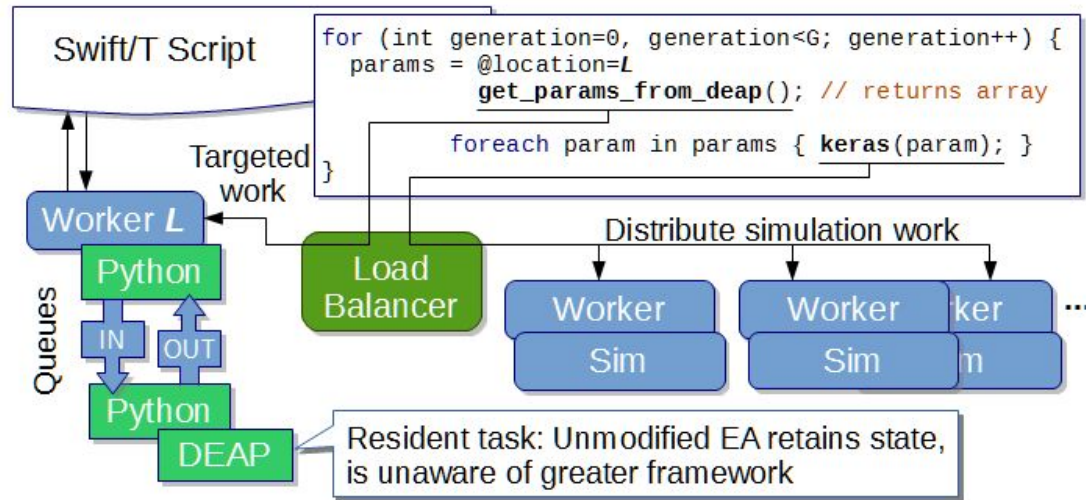


- The core novel contributions of EMEWS are shown in green, these allow the Swift script to access a running **Model Exploration (ME)** algorithm, and create an **inversion of control (IoC)** workflow
- Both green and blue boxes accept **existing multi-language code**
- <http://emews.org>



# EMEWS: Extreme-scale model exploration workflows in Swift/T

- How do we couple complex model exploration algorithms with workflows?  
Optimization, active learning, uncertainty quantification...



- From desktop to large-scale model exploration with Swift/T  
Proc. Winter Simulation Conference 2016

# Links

- Swift/T Home: <http://swift-lang.org/Swift-T>
- Swift/T Guide: <http://swift-lang.github.io/swift-t/guide.html>
- Swift/T Sites Guide: <http://swift-lang.github.io/swift-t/sites.html>
- Swift/T GitHub: <https://github.com/swift-lang/swift-t>
- Support: <https://groups.google.com/forum/#!forum/swift-t-user>
- Book chapter (easiest introduction): [http://www.mcs.anl.gov/~wozniak/papers/ProgrammingModels\\_Swift\\_2015.pdf](http://www.mcs.anl.gov/~wozniak/papers/ProgrammingModels_Swift_2015.pdf)
- Other papers: <http://swift-lang.github.io/swift-t/pubs.html>