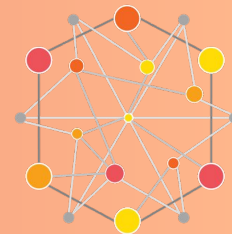


Tutorial: Developing Robust and Scalable Next Generation Workflows Applications and Systems

ISC-HPC 2022





`https://parsl-project.org/`

Parsl: a parallel programming library for Python

Apps define opportunities for **parallelism**
Python apps call Python functions
Bash apps call external applications

Apps return “futures”: a proxy for a result that might not yet be available

Apps run concurrently respecting dataflow dependencies. Natural parallel programming!

Parsl scripts are independent of where they run. Write once run anywhere!

```
pip install parsl
```

```
@python_app
def hello():
    return 'Hello World!'

print(hello().result())
```



Hello World!

```
@bash_app
def echo_hello(stdout='echo-hello.stdout'):
    return 'echo "Hello World!"'

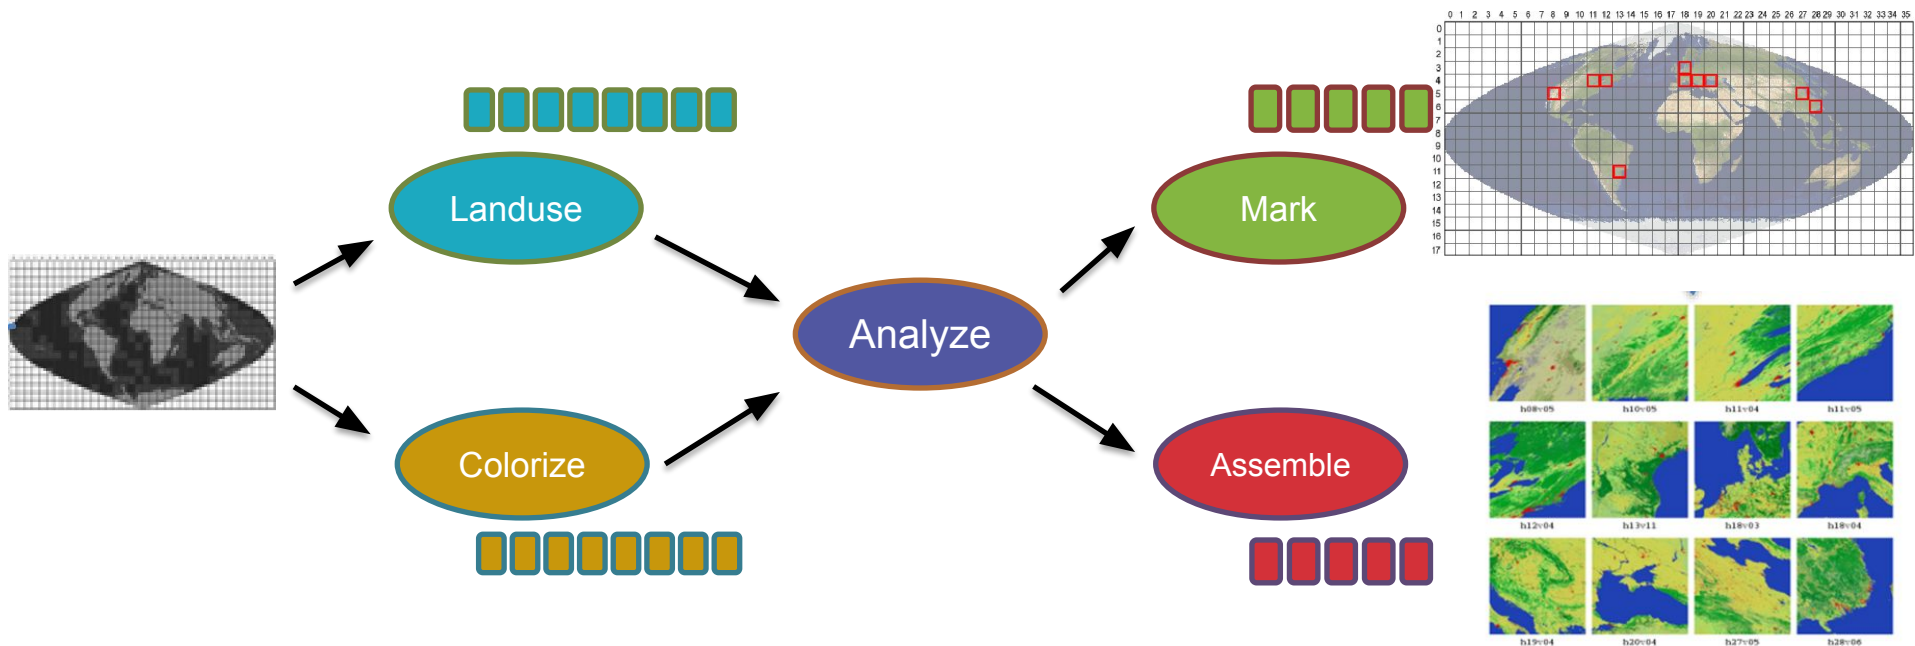
echo_hello().result()

with open('echo-hello.stdout', 'r') as f:
    print(f.read())
```



Hello World!

Data-driven example: parallel geospatial analysis



Land-use Image processing pipeline for the MODIS remote sensor

Expressing parallelism using Parsl

1) *Wrap the science applications as Parsl Apps:*

```
@bash_app
def landuse(img, outputs=[]):
    return './landuse_sim.sh {} {}'.format(img, outputs[0])

@python_app
def colorize(img, outputs=[]):
    return color_package(img, len(outputs))

@python_app
def analyze(land_chunks, color_chunks):
    return combine(land_chunks, color_chunks)
```

Expressing a many task workflow in Parsl

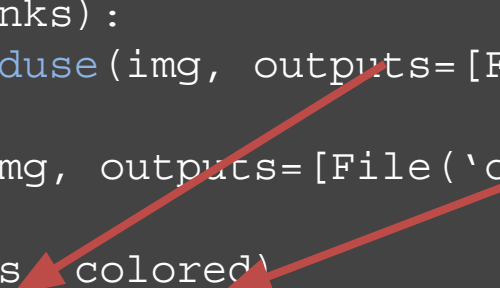
2) *Execute the parallel workflow by calling Apps:*

```
lchunks = []

for i in range (nchunks):
    lchunks.append(landuse(img, outputs=[File('lc-%s.txt' % i)]))

colored = colorize(img, outputs=[File('c-%s' % i) for i in range(5)])

all = analyze(lchunks, colored)
```



Decomposing dynamic parallel execution into a task-dependency graph

jupyter parsl-introduction (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Monte Carlo workflow

Many scientific applications use the [monte-carlo method](#) to compute results.

If a circle with radius r is inscribed inside a square with side length $2r$ then the area of the circle is πr^2 and the area of the square is $(2r)^2$. Thus, if N uniformly distributed random points are dropped within the square then approximately $N\pi/4$ will be inside the circle.

Each call to the function `pi()` is executed independently and in parallel. The `avg_three()` app is used to compute the average of the futures that were returned from the `pi()` calls.

The dependency chain looks like this:

```
App Calls    pi() pi() pi()
              \  |  /
Futures      a  b  c
              \  |  /
App Call     avg_points()
              |
Future       avg_pi
```

In []:

```
# App that estimates pi by placing points in a box
@python_app
def pi(total):
    import random

    # Set the size of the box (edge length) in which we drop random points
    edge_length = 10000
    center = edge_length / 2
    c2 = center ** 2
    count = 0

    for i in range(total):
        # Drop a random point in the box.
        x,y = random.randint(1, edge_length), random.randint(1, edge_length)
        # Count points within the circle
        if (x-center)**2 + (y-center)**2 < c2:
            count += 1


    return (count*4/total)

# App that computes the average of the values
@python_app
def avg_points(a, b, c):
    return (a + b + c)/3

# Estimate three values for pi
a, b, c = pi(10**6), pi(10**6), pi(10**6)

# Compute the average of the three estimates
avg_pi = avg_points(a, b, c)

# Print the results
print("A: {0:.5f} B: {1:.5f} C: {2:.5f}".format(a.result(), b.result(), c.result()))
print("Average: {0:.5f}".format(avg_pi.result()))
```



Parsl programs can be executed in different ways on different systems




Executors (concurrent.futures.Executor interface)

HTEX



Work Queue



Flux



EXEX



RADICAL-Cybertools



funcX



IPyParallel

IP[y]:

Production

Prototype

Deprecated

Providers

Slurm

LSF

GridEngine

Kubernetes

AWS

PBS

Cobalt

HTCondor

Google

Ad hoc

Parsl executors scale to 2M tasks/256K workers

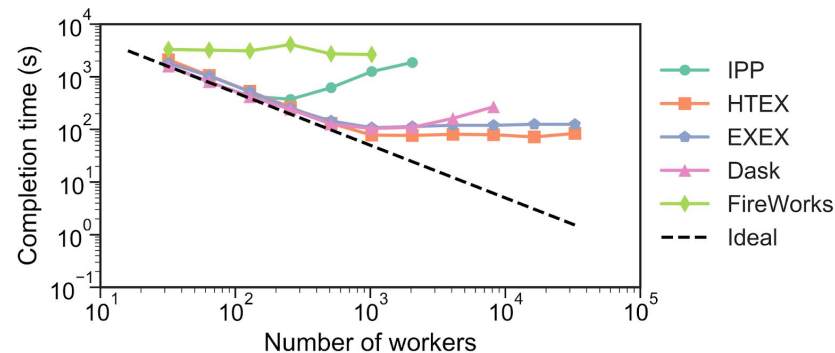
HTEX and EXEX outperform other Python-based approaches

Parsl scales to more than 250K workers (8K nodes) and ~2M tasks


Framework	Maximum # of workers [†]	Maximum # of nodes [†]	Maximum tasks/second [‡]
Parsl-IPP	2048	64	330
Parsl-HTEX	65 536	2048 [*]	1181
Parsl-EXEX	262 144	8192 [*]	1176
FireWorks	1024	32	4
Dask distributed	4096	128	2617

Babuji et.al. "Parsl: Pervasive Parallel Programming in Python."
ACM International Symposium on High-Performance Parallel and
Distributed Computing (HPDC). 2019.

Strong scaling (50K 1s tasks)





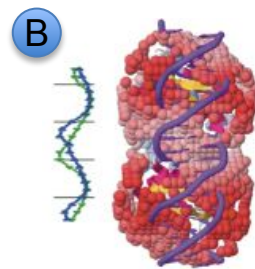
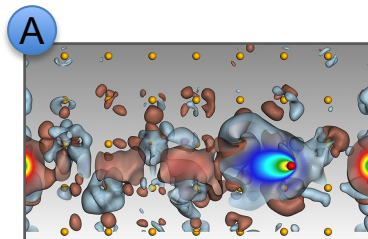


Weak scaling (10 1s tasks per worker)



Parsl is being used in a wide range of scientific applications

- A Machine learning to predict stopping power in materials
- B Protein and biomolecule structure and interaction
- C LSST simulation and weak lensing using sky surveys
- D Cosmic ray showers in QuarkNet
- E Information extraction to classify image types in papers
- F Materials science at the Advanced Photon Source
- G Machine learning and data analytics in materials



More examples:

<https://parsl-project.org/parslfest>

Other functionality provided by Parsl



Resource abstraction

Block-based model overlaying different providers and resources



Fault tolerance

Support for retries, checkpointing, and memoization



Multi site

Combining executors/providers for execution across different resources



Elasticity

Automated resource expansion/retraction based on workload



Monitoring

Workflow and resource monitoring and visualization



Globus

Delegated authentication and wide area data management



Data management

Automated staging with HTTP, FTP, and Globus



Containers

Sandboxed execution environments for workers and tasks



Jupyter integration

Seamless description and management of workflows



Reproducibility

Capture of workflow provenance in the task graph

Exercises

`https://github.com/ExaWorks/Tutorial`